



BSc (Hons) in **Information Technology**
Specialized in AI

IT3012 : Intelligent Agents

Year 3 · Semester 1 · 2026 Semester 2

Faculty of Computing



Practical 04

Objective & Lecture Mapping: In previous labs, you explored Uninformed Search strategies like BFS and UCS inside `agent.py` [cite: 1]. Today, we transition to **Informed Search**. You will enhance your agent by implementing the **A* Search** algorithm, using **Heuristics** to drastically reduce the number of explored nodes in the `visual_grid_game.py` environment [cite: 4]. You will start with basic heuristic functions and connect them to your search logic.

Part 1: Practical Implementation — Building the A* Agent

Step 1.1: Implementing the Heuristic Functions

Informed search algorithms require a heuristic function, $h(n)$, which estimates the cheapest path from the current node to the goal. You will calculate distance metrics on a 2D grid.

1. Create a new Branch called `Week_04` in your Forked Github Repo.
2. Open `agent.py` and locate your `SearchAgent` class [cite: 1].
3. Add a new method named `def manhattan_distance(self, pos, goal):` that calculates the distance using the formula $h(n) = |x_1 - x_2| + |y_1 - y_2|$ and returns the integer value.
4. Add a second method named `def euclidean_distance(self, pos, goal):` that calculates the straight-line distance using the formula $h(n) = \sqrt{(x_1 - x_2)^2 + (y_1 - y_2)^2}$. You will need to `import math` for this.
5. *Testing Checkpoint:* Print the outputs of both functions for a mock start position (0, 0) and goal (3, 4) to verify that Manhattan returns 7 and Euclidean returns 5.0.

Step 1.2: Implementing A* Search

A* evaluates nodes by combining the path cost so far, $g(n)$, and the estimated cost to the goal, $h(n)$. The evaluation function is $f(n) = g(n) + h(n)$.

1. Inside `SearchAgent`, create a new method: `def astar_search(self, start_pos, goal_pos, walls, grid_size, heuristic_type='manhattan')`.
2. Initialize an empty priority queue using `heapq` and an empty set for `reached_states`.
3. Push the initial state into the priority queue. Unlike UCS which only uses $g(n)$, your A* tuple must be formatted as: `(f_cost, g_cost, current_pos, path_taken)`. For the starting node, $g(n) = 0$. Calculate $h(n)$ using your chosen heuristic method, and set $f(n) = g(n) + h(n)$.
4. Create your standard `while` loop to process the queue. Pop the node with the lowest `f_cost`. If `current_pos == goal_pos`, return the `path_taken`. Otherwise, add `current_pos` to `reached_states`.
5. For node expansion, check the four adjacent cells (Up, Down, Left, Right). For each valid neighbor (not a wall, within bounds, and not reached): Calculate $g_{\text{new}} = g_{\text{current}} + 1$, h_{new} using your heuristic, and $f_{\text{new}} = g_{\text{new}} + h_{\text{new}}$. Push the new tuple to the priority queue.

Step 1.3: Integrating A* into the Agent's Decision Loop

Finally, you need to connect your new search algorithm to the agent's perception and action loop so it can run in the visual environment.

1. Navigate to the `sense_and_act(self, percept)` method in your `SearchAgent`.
2. Add an `elif self.active_algo == 'AStar':` block to handle the new algorithm alongside your existing BFS/DFS/UCS.
3. Extract the necessary global state from the percept dictionary (e.g., `percept['remaining_food']`).
4. Find the closest food item to act as the `goal_pos`. Call your `astar_search` method and store the returned list of actions in `self.plan`.

5. Open `visual_grid_game.py` and modify the `GridGameGUI` initialization to inject your `SearchAgent()` instead of the `ModelBasedAgent()`. Run the simulation and observe the optimized pathfinding!

Part 2: Theoretical Evaluation

Based on your implementation and Lecture 05, complete the following analytical questions.

1. (Understand) What is the key difference between how Uniform-Cost Search (UCS) and A* Search prioritize which node to explore next?

UCS prioritizes nodes using only the cost already travelled:

$$f(n) = g(n)$$

In `agent.py:163-181`, the priority queue selects the node with the smallest accumulated path cost.

A* combines the travelled cost with an estimated remaining cost:

$$f(n) = g(n) + h(n)$$

Therefore, A* prefers nodes that are both cheap to reach and estimated to be close to the goal. UCS has no goal-directed guidance, while A* uses the heuristic.

2. (Analyze) In Step 1.1, you used Manhattan Distance. Why is Manhattan Distance considered an "admissible" heuristic for this specific 4-way movement grid, and what would happen to your A* algorithm if the heuristic was NOT admissible?

This grid allows movement only up, down, left, or right, with each move costing 1. Manhattan distance is:

$$h(n) = |x_n - x_g| + |y_n - y_g|$$

It counts the minimum number of horizontal and vertical moves needed if there

were no walls. Walls can only force extra movement; they cannot create a route shorter than this direct minimum.

Therefore:

$$h(n) \leq h^*(n)$$

where $h^*(n)$ is the actual cheapest remaining path cost. This makes Manhattan distance admissible.

If the heuristic overestimated the true cost, A^* could expand a more expensive-looking node later even though it belongs to the optimal path. In that case, A^* would no longer be guaranteed to return an optimal path.

3. (Evaluate) If we modified `visual_grid_game.py` to allow the agent to move diagonally (8-way movement), would Manhattan distance still be an admissible heuristic? Why or why not? Which metric should you switch to?

With 8-way movement and equal cost for horizontal, vertical, and diagonal moves, Manhattan distance is no longer admissible.

For example, from (0, 0) to (3, 3):

Manhattan distance: $3 + 3 = 6$

Actual diagonal path: 3 moves

Manhattan distance overestimates the true cost, violating admissibility.

For equal-cost diagonal movement, use Chebyshev distance:

$$h(n) = \max(|x_n - x_g|, |y_n - y_g|)$$

If diagonal movement has a different cost, use an appropriate octile distance heuristic instead.

4. (Create) When targeting multiple food items simultaneously, calculating the distance to just the single closest food item is a weak heuristic. Propose (in text) a stronger heuristic for navigating the grid to eat ALL remaining food efficiently.

The current agent plans to only one food item at a time, selecting the nearest pellet. This can be inefficient because the locally closest pellet may leave the remaining pellets far apart.

A stronger approach is to treat the state as:

(agent position, set of remaining food)

A useful heuristic could be:

$h(n)$ = distance from agent to the nearest remaining food + MST cost of all remaining food

where MST is the minimum spanning tree connecting all remaining food positions using grid distances.

This estimates:

- The cost to reach the food region.
- The minimum cost to connect and visit all remaining food.

For a simpler but weaker version, calculate the minimum distance from the agent to one food plus the minimum spanning tree cost between all pellets. This is generally stronger than choosing only the closest pellet, while remaining a lower-bound estimate when based on obstacle-aware shortest-path distances.

Once Completed Get a PDF of the completed File and Push into the Week 04 branch in the Github Project

Finalize and Lock Lab 04 Documentation

